

### 361 HW3 Report

It is quite easy to see that any program that behaves like infinite memory is desirable. What does this really mean? Any program with the fewest number of page misses and highest number of page hits is desirable. Based on the sample input files given, it is clear that photographs with little to no variation in color value would perform the best. But having small variation in color value is not enough to succeed. The timing of the variation is also important. If there was just enough variation in color value to exceed the the memory size AND this variation was spread out in very quick succession, it is obvious that the cache would perform poorly. Take the two pictures supplied that performed poorly, `light_drops.raw` and `GrandeJatte.raw`. If we look at the pictures themselves, we can see there is a lot of **unsustained** variation. "Reading" from left to right, we are constantly seeing new colors (colors unlike the previous 256 seen) In the case of `light_drops` it is probably the same few colors per line. In `GrandeJatte` we are seeing all sorts of colors. In either case, there is a poor locality of reference as well as poor sustained referencing. Contrasting this, if we look at `VRUPL_logo.raw` versus `LUG_newbieews.raw`, we can see the same trends but this time with high locality of reference. The locality of reference is evident by the low variation of colors. Again, we see that variation in color is not enough. Here, we can see that `VRUPL`'s diagonal content aids in seeing the same color for shorts bursts of time, which is performant for our page table.

What the data says:

As we analyze based on the size of the frame table, I will start with a small frame size and increase until the pictures perform near unlimited. From the start, we can see that `VRUPL` is far more performant than the other three. `LUG` performs at 50% of its max, `light_drops` at 18%, and `GrandeJatte` at 11%. If we take a look at average number of times swapped, we see something quite interesting. We see that `LUG` has a larger number of times swapped yet more hit percentage than `light_drops`. This is most likely from the fact that `LUG` has very few **types** of colors but very frequent **changes**. `light_drops` has many types (shades) of colors, so the average is actually very low (each very specific color may only be needed a few times).

Let us look at the difference between a frame table size of 1536 and 3072.

1536: `light_drops` percent of max: 58% `LUG`: 72%

3072: `light_drops`: 81% `LUG`: 78%

Why would the overall performance of `light_drops` somehow radically surpass that of `LUG`?

I would attribute this to `LUG`'s **change** in average times switched. Basically, how helpful was the increase in the cache in terms of times of replacement. For `light_drops`, the average is halved from 1536 to 3072! Yet for `LUG`, we are left with 80% of the previous run's average number of times swapped. It must be the case that despite `LUG`'s simplicity, the cache is still not enough to capture small variations in color from letter to letter. What does this mean in terms of frame

tables? Gradual changes to required page tables are preferred to sharp drastic ones. This seems fairly obvious but the data brings us to this conclusion logically.